

CS231n- Lecture 1

February 11, 2017

1 Image Classification

Take image and classify into one label

Semantic Gap: 300 x 100 x 3 being the number of RGB number and size

Problems:

- Work with millions of numbers is hard.

- Intrinsics of Cameras are different

- Illumination affects images

- Deformation affects images

- Occlusion (can't see full objects like cat)

- Background color

- \sum

- Intra Class Variation

Rule based approaches are super hard because it is not scalable.

Data has tremendously increased \Rightarrow better training

First Classifier:

- Nearest Neighbor Classifier

- CIFAR-10

- 10 labels

- 50k training images, 32x32

- 10k test images

How can we compare Images: Distance metric

- Manhattan Distance(L1)

- Return label with $\text{argmin}(Md)$

- Linearly slower as the training set is bigger.

CNNs work much faster after training is done(const)

We can use L2(Euclidean Dist)

This distance is a hyperparameter.

If we generalise it to the k-Nearest Neighbors we'll get k nearest neighbors and then vote

As k increases it tends to smooth out.

Choice of k is a hyperparameter too

Nearest neighbor gives 100% accuracy on training set when using Euclidean norm or Manhattan Norm

K-nearest neighbor doesn't give a 100% accuracy on training set
 Hyperparameter testing can be done with trial and errors
 Split training and test data to avoid such bs.
 We separate Training data sometimes to 5 folds and then train on 4 folds
 and test on the other one(say 5th fold to tune hypparams) This fold 5 is cal
 called Validation set Cross Validation is iterating accross these folds as
 validation sets and then average for each.
 We find k=7 to be best as it is peeking a plot of acc vs k
 k-NN is never used: Terrible performance at testing time
 Distance metrics on level of whole images can be unintuitive
 Linear Classification:
 NNs can see, hear, language, control(robots)
 Building blocks involved.
 Image Captioning:
 CNN(to see) -> RNN (to model sequences(in this case words))
 KNN is nonparametric
 in Lin Classification: $f(x, W)$ x-image, W is parameters to give 10 numbers
 indicating class scores
 Image [32x32x3] we stretch it into a long vector
 Suppose $f(x, W) = W \begin{matrix} x \\ 10 \times 1 \end{matrix} + b \begin{matrix} 3072 \times 1 \\ 10 \times 1 \end{matrix}$
 We resize images in this approach to 32x32 although not optimal
 Score is a weighted sum of all the pixel vals
 if you reshape rows of W to image we see a weirdish template image of each
 label
 Mostly only captures colors which means it is not good at all
 GreyScale will work terribly in linear classifiers too
 The score is the output we get from this fn